

PARPING

PROJET GALAX

Florian Huynh, Tchadel Icard

Introduction

Galax est une simulation de collision entre deux galaxies basée sur le problème des N-corps

- Objectif : maximiser le nombre de FPS pour 10,000 particules sans affichage
- Techniques : OpenMD, SIMD (AVX2), FMA

Implémentation naïve

- Complexité algorithmique : $O(n^2)$
- Objectif : utiliser cette implémentation comme base de comparaison
- Résultats
 - 2000 particules → 50 FPS
 - 5000 particules → 8.4 FPS
 - 10000 particules → 2 FPS

```

11 void Model_CPU_naive
12 ::step()
13 {
14     std::fill(accelerationsx.begin(), accelerationsx.end(), 0);
15     std::fill(accelerationsy.begin(), accelerationsy.end(), 0);
16     std::fill(accelerationsz.begin(), accelerationsz.end(), 0);
17
18     for (int i = 0; i < n_particles; i++)
19     {
20         for [int j = 0; j < n_particles; j++]
21         {
22             if(i != j)
23             {
24                 const float diffx = particles.x[j] - particles.x[i];
25                 const float diffy = particles.y[j] - particles.y[i];
26                 const float diffz = particles.z[j] - particles.z[i];
27
28                 float dij = diffx * diffx + diffy * diffy + diffz * diffz;
29
30                 if (dij < 1.0)
31                 {
32                     dij = 10.0;
33                 }
34                 else
35                 {
36                     dij = std::sqrt(dij);
37                     dij = 10.0 / (dij * dij * dij);
38                 }
39
40                 accelerationsx[i] += diffx * dij * initstate.masses[j];
41                 accelerationsy[i] += diffy * dij * initstate.masses[j];
42                 accelerationsz[i] += diffz * dij * initstate.masses[j];
43             }
44         }
45     }
46
47     for (int i = 0; i < n_particles; i++)
48     {
49         velocitiesx[i] += accelerationsx[i] * 2.0f;
50         velocitiesy[i] += accelerationsy[i] * 2.0f;
51         velocitiesz[i] += accelerationsz[i] * 2.0f;
52         particles.x[i] += velocitiesx [i] * 0.1f;
53         particles.y[i] += velocitiesy [i] * 0.1f;
54         particles.z[i] += velocitiesz [i] * 0.1f;
55     }
56 }
```

Mathieu Léonardon, 4 years ago • Initial code

Implémentation CPU Fast avec OpenMP

- Utilisation d'OpenMP pour pouvoir faire tourner l'algorithme de manière parallèle (4 coeurs de la machine)
- Accélération attendue de 4

```

23     void Model_CPU_fast
24     ::step()
25     {
26         std::fill(accelerationsx.begin(), accelerationsx.end(), 0);
27         std::fill(accelerationsy.begin(), accelerationsy.end(), 0);
28         std::fill(accelerationsz.begin(), accelerationsz.end(), 0);
29
30         #pragma omp parallel for schedule(static)
31         for (int i = 0; i < n_particles; i++)
32         {
33             for (int j = 0; j < n_particles; j++)
34             {
35                 if(i != j)
36                 {
37                     const float diffx = particles.x[j] - particles.x[i];
38                     const float diffy = particles.y[j] - particles.y[i];
39                     const float diffz = particles.z[j] - particles.z[i];
40
41                     float dij = diffx * diffx + diffy * diffy + diffz * diffz;
42
43                     if (dij < 1.0)
44                     {
45                         dij = 10.0;
46                     }
47                     else
48                     {
49                         dij = std::sqrt(dij);
50                         dij = 10.0 / (dij * dij * dij);
51                     }
52
53                     accelerationsx[i] += diffx * dij * initstate.masses[j];
54                     accelerationsy[i] += diffy * dij * initstate.masses[j];
55                     accelerationsz[i] += diffz * dij * initstate.masses[j];
56                 }
57             }
58         }
59     }

```

Implémentation CPU Fast avec OpenMP

CPU	FPS	Particules	Accélération
CPU Naïve	50	2000	REF
CPU Naïve	8.4	5000	REF
CPU Naïve	2	10000	REF
CPU Fast (OpenMP)	150	2000	3
CPU Fast (OpenMP)	20	5000	2,38
CPU Fast (OpenMP)	5,6	10000	2,8

Résultat inférieur à l'accélération attendue

Implémentation CPU Fast avec OpenMP + SIMD v1

- Utilisation de SIMD pour exécuter une instruction sur plusieurs données en parallèle, accélérant les calculs.
- Accélération attendue de 8

```
#pragma omp parallel for schedule(static)
for (int i = 0; i < n_particles; i++)
{
    b_type acc_x = b_type(0.0f);
    b_type acc_y = b_type(0.0f);
    b_type acc_z = b_type(0.0f);
    for (int j = 0; j < n_particles; j += b_type::size)
    {
        if(i != j)
        {
            b_type xj = b_type::load_unaligned(&particles.x[j]);
            b_type yj = b_type::load_unaligned(&particles.y[j]);
            b_type zj = b_type::load_unaligned(&particles.z[j]);
            b_type mj = b_type::load_unaligned(&initstate.masses[j]);

            b_type xi = b_type(particles.x[i]);
            b_type yi = b_type(particles.y[i]);
            b_type zi = b_type(particles.z[i]);

            b_type diffx = xj - xi;
            b_type diffy = yj - yi;
            b_type diffz = zj - zi;

            b_type dij = diffx * diffx + diffy * diffy + diffz * diffz;

            // Compute inverse squared-root and handle singularities
            auto mask = dij < b_type(1.0f);
            dij = xs::select(mask, b_type(10.0f), xs::rsqrt(dij) * xs::rsqrt(dij) * b_type(10.0f));

            acc_x += diffx * dij * mj;
            acc_y += diffy * dij * mj;
            acc_z += diffz * dij * mj;
        }
    }
    // Store results back
    accelerationsx[i] += xs::reduce_add(acc_x);
    accelerationsy[i] += xs::reduce_add(acc_y);
    accelerationsz[i] += xs::reduce_add(acc_z);
}
```

Implémentation CPU Fast avec OpenMP + SIMD V2

- Nouvelle version avec les constantes qui sont définies en haut
- Remplacement du select par un max (moins couteux en calculs)
- On calcule qu'une seule fois la racine inverse au lieu de 3 fois avant

```

227 const b_type b10 = b_type(10.0f);
228 const b_type b1 = b_type(1.0f);
229
230 #pragma omp parallel for schedule(static)
231 for (int i = 0; i < n_particles; i++)
232 {
233     b_type acc_x = b_type(0.0f);
234     b_type acc_y = b_type(0.0f);
235     b_type acc_z = b_type(0.0f);
236
237     b_type xi = b_type(particles.x[i]);
238     b_type yi = b_type(particles.y[i]);
239     b_type zi = b_type(particles.z[i]);
240
241     for (int j = 0; j < n_particles; j += b_type::size)
242     {
243         if (i != j) {
244             b_type xj = b_type::load_unaligned(&particles.x[j]);
245             b_type yj = b_type::load_unaligned(&particles.y[j]);
246             b_type zj = b_type::load_unaligned(&particles.z[j]);
247             b_type mj = b_type::load_unaligned(&initstate.masses[j]);
248
249             b_type diffx = xj - xi;
250             b_type diffy = yj - yi;
251             b_type diffz = zj - zi;
252
253             b_type dij = diffx * diffx + diffy * diffy + diffz * diffz;
254             dij = xs::max(dij, b1);
255
256             // Compute inverse squared-root and handle singularities
257             b_type inv_dij = xs::rsqrt(dij);
258             inv_dij = inv_dij * inv_dij * inv_dij;
259
260             dij = b10 * inv_dij;
261
262             acc_x += diffx * dij * mj;
263             acc_y += diffy * dij * mj;
264             acc_z += diffz * dij * mj;
265
266         }
267
268     // Store results back
269     accelerationsx[i] += xs::reduce_add(acc_x);
270     accelerationsy[i] += xs::reduce_add(acc_y);
271     accelerationsz[i] += xs::reduce_add(acc_z);
272 }

```

Implémentation CPU Fast avec OpenMP + SIMD V3

- On charge 8 particules à la fois et on calcule toutes les interactions avec un rotate
- Permet d'optimiser les accès mémoire
- Elle ne réduit pas le nombre d'opérations effectuées

```

418 const b_type b10 = b_type(10.0f);
419 const b_type b1 = b_type(1.0f);
420
421 #pragma omp parallel for
422 for (int i = 0; i < n_particles; i += b_type::size)
423 {
424     b_type xi = b_type::load_unaligned(&particles.x[i]);
425     b_type yi = b_type::load_unaligned(&particles.y[i]);
426     b_type zi = b_type::load_unaligned(&particles.z[i]);
427
428     b_type acc_x = b_type(0.0f);
429     b_type acc_y = b_type(0.0f);
430     b_type acc_z = b_type(0.0f);
431     for (int j = 0; j < n_particles; j += b_type::size)
432     {
433         b_type xj = b_type::load_unaligned(&particles.x[j]);
434         b_type yj = b_type::load_unaligned(&particles.y[j]);
435         b_type zj = b_type::load_unaligned(&particles.z[j]);
436         b_type mj = b_type::load_unaligned(&initstate.masses[j]);
437
438         for (int k = 0; k < b_type::size; k++)
439         {
440             b_type diffx = xj - xi;
441             b_type diffy = yj - yi;
442             b_type diffz = zj - zi;
443
444             b_type dij = diffx * diffx + diffy * diffy + diffz * diffz;
445             dij = xs::max(dij, b1);
446
447             // Compute inverse squared-root and handle singularities
448             b_type inv_dij = xs::rsqrt(dij);
449             inv_dij = inv_dij * inv_dij * inv_dij;
450
451             dij = b10 * inv_dij;
452
453             acc_x += diffx * dij * mj;
454             acc_y += diffy * dij * mj;
455             acc_z += diffz * dij * mj;
456
457             xj = xs::rotate_right<1>(xj);
458             yj = xs::rotate_right<1>(yj);
459             zj = xs::rotate_right<1>(zj);
460             mj = xs::rotate_right<1>(mj);
461         }
462
463         acc_x.store_unaligned(&accelerationsx[i]);
464         acc_y.store_unaligned(&accelerationsy[i]);
465         acc_z.store_unaligned(&accelerationsz[i]);
466     }
}

```

Implémentation CPU Fast avec OpenMP

CPU	FPS	Particules	Accélération
CPU Naïve	2	10000	REF
CPU Fast V0	5,6	10000	2,8 (CPU_Naive)
CPU Fast V1	65	10000	32,5 (CPU_Naive) 11,6 (CPU_Fast)
CPU Fast V2	84	10000	42 (CPU_Naive) 15 (CPU_Fast)
CPU Fast V3	90 (110 sans GNOME)	10000	45 (CPU_Naive) 16 (CPU_Fast)

Accélération attendue liée à SIMD de 8

Optimisation effectuée par le compilateur ?

Mesures avec OpenMP pour 10 000 particules

- OpenMP avec 1 thread : ~1,5 FPS
- OpenMP avec 4 threads : ~5,2-6 FPS

Pour comparer avec la version naïve

- CPU Naïve avec 10 000 particules : ~2,1 FPS

On a bien une accélération de 4 avec le même code

Validate

```
f22huynh@fl-tp-br-633:~/galax_eleves/build$ ./bin/galax -c CPU_FAST -n 10000 --display NO --validate
State updates per second: 0 ; average distance vs reference: 1.94313e-05; min error : 0; max error : 0.000111888
State updates per second: 0 ; average distance vs reference: 5.71617e-05; min error : 0; max error : 0.000336332
State updates per second: 0 ; average distance vs reference: 9.86522e-05; min error : 0; max error : 0.000580332
State updates per second: 0 ; average distance vs reference: 0.000113735; min error : 0; max error : 0.000528291
State updates per second: 0 ; average distance vs reference: 0.000134336; min error : 0; max error : 0.00066542
State updates per second: 0 ; average distance vs reference: 0.000183676; min error : 0; max error : 0.000979989
State updates per second: 0 ; average distance vs reference: 0.000222196; min error : 0; max error : 0.00108683
State updates per second: 0 ; average distance vs reference: 0.000249959; min error : 0; max error : 0.00111195
State updates per second: 0 ; average distance vs reference: 0.00029546; min error : 0; max error : 0.00144733
State updates per second: 124.59 ; average distance vs reference: 0.00034518; min error : 0; max error : 0.0017722
State updates per second: 124.59 ; average distance vs reference: 0.00039891; min error : 0; max error : 0.00286665
State updates per second: 124.59 ; average distance vs reference: 0.000469698; min error : 0; max error : 0.00267574
State updates per second: 124.59 ; average distance vs reference: 0.000533009; min error : 0; max error : 0.00310962
State updates per second: 124.59 ; average distance vs reference: 0.00054639; min error : 0; max error : 0.00370819
State updates per second: 124.59 ; average distance vs reference: 0.000597462; min error : 0; max error : 0.00494501
State updates per second: 124.59 ; average distance vs reference: 0.000700545; min error : 0; max error : 0.00461178
State updates per second: 124.59 ; average distance vs reference: 0.000790951; min error : 0; max error : 0.00608116
State updates per second: 124.59 ; average distance vs reference: 0.000841109; min error : 0; max error : 0.00895351
State updates per second: 124.59 ; average distance vs reference: 0.000908443; min error : 0; max error : 0.010299
State updates per second: 124.59 ; average distance vs reference: 0.00103664; min error : 0; max error : 0.013408
State updates per second: 124.59 ; average distance vs reference: 0.00110238; min error : 0; max error : 0.0211925
State updates per second: 124.59 ; average distance vs reference: 0.00115024; min error : 0; max error : 0.0136486
State updates per second: 124.59 ; average distance vs reference: 0.00131057; min error : 0; max error : 0.013997
State updates per second: 124.59 ; average distance vs reference: 0.00147942; min error : 0; max error : 0.018734
State updates per second: 124.59 ; average distance vs reference: 0.00155534; min error : 0; max error : 0.0192557
State updates per second: 124.59 ; average distance vs reference: 0.00162146; min error : 0; max error : 0.02438
State updates per second: 124.59 ; average distance vs reference: 0.00179331; min error : 9.53674e-07; max error : 0.0350508
State updates per second: 124.59 ; average distance vs reference: 0.00192743; min error : 4.76837e-07; max error : 0.047497
^CState updates per second: 124.59 ;
State updates per second: 61.977 ; average distance vs reference: 0.00208586; min error : 4.76837e-07; max error : 0.0624085
average distance vs reference: 0.00231579; min error : 9.53674e-07; max error : 0.080265
```

MERCI

